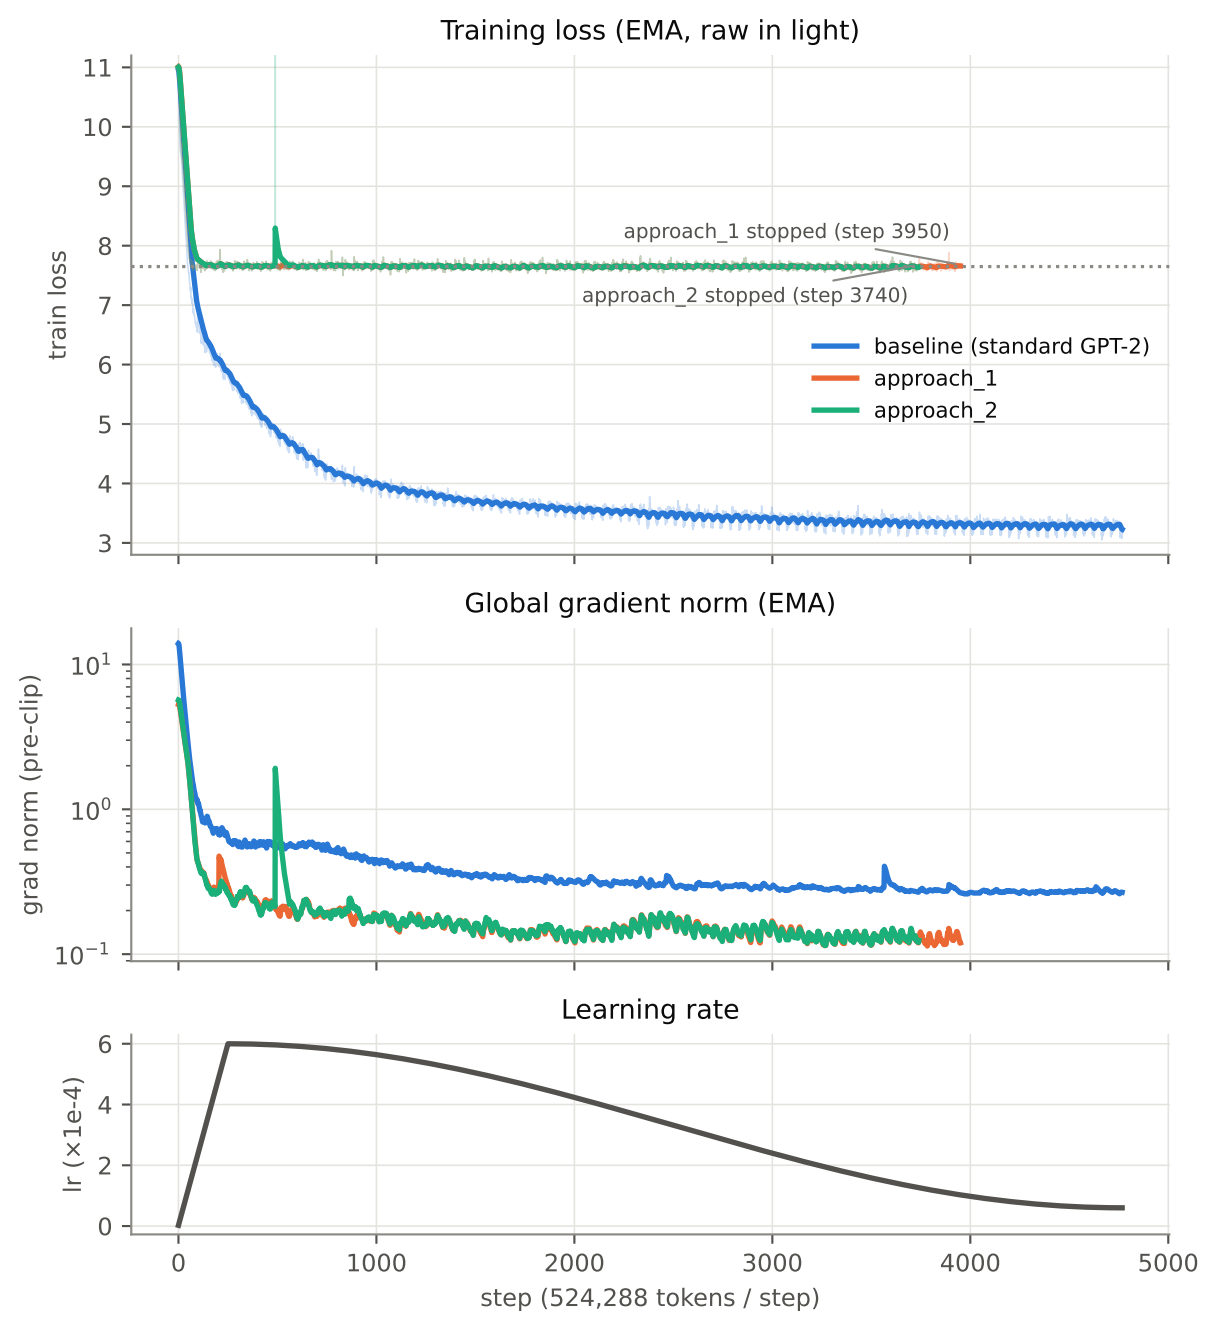

Training loss (EMA, raw in light)
train loss
approach_1 stopped (step 3950)
approach_2 stopped (step 3740)
baseline (standard GPT-2)
approach_1
approach_2
Global gradient norm (EMA)
grad norm (pre-clip)
Learning rate
lr (×1e-4)
step (524,288 tokens / step)