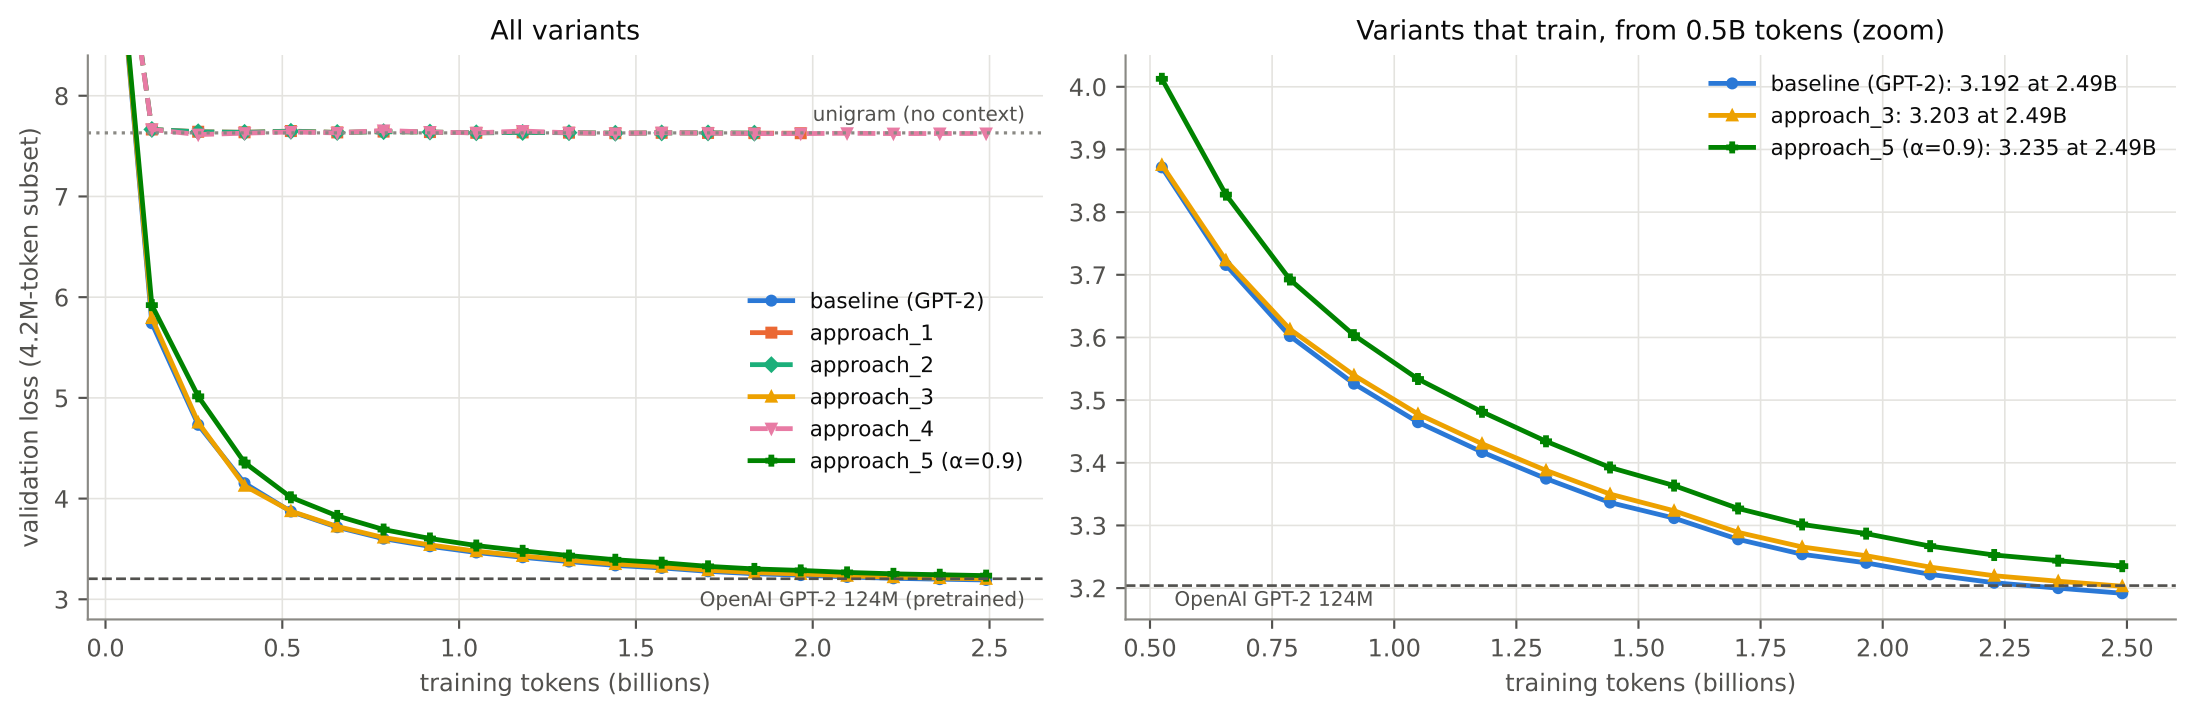
All variants
validation loss (4.2M-token subset)
training tokens (billions)
unigram (no context)
OpenAI GPT-2 124M (pretrained)
baseline (GPT-2)
approach_1
approach_2
approach_3
approach_4
approach_5 (α=0.9)
Variants that train, from 0.5B tokens (zoom)
training tokens (billions)
OpenAI GPT-2 124M
baseline (GPT-2): 3.192 at 2.49B
approach_3: 3.203 at 2.49B
approach_5 (α=0.9): 3.235 at 2.49B